



SD TSIA 204
Linear Models
Intro to linear models

Ekhiñe Irurozki
Télécom Paris


A 2D starting example



A 2D starting example



A 2D starting example



Notation interpretation

- ▶ $n = 244$
- ▶ $p = 1$
- ▶ y_i : tip let by the i -th customer
- ▶ x_i : total bill paid by the i -th customer
- ▶ y : the observation is the tips, dependent variable
- ▶ x : the feature/covariate, price of the bill, independent variable

Linear model / Linear regression hypothesis : assume that the price of the bill and the tip let are linearly correlated

Exo : use `describe()` from `Pandas` to get a rough data summary

Three questions to be covered : modeling, learning and predicting

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
# Generate example data
np.random.seed(42)
X = np.random.rand(20, 1)*10 # Independent variable
y = 2 * X + 3 + np.random.randn(20, 1) # Dependent variable
# Fit linear regression model
model = LinearRegression()
model.fit(X, y)
# Predict y values using the model
X_new = np.linspace(0, 10, 100).reshape(-1, 1)
y_pred = model.predict(X_new)
# Create a scatter plot of the data points
plt.scatter(X, y, label='Data Points')
# Plot the linear regression line
plt.plot(X_new, y_pred, color='red', label='Linear Regression Line')
plt.xlabel('X')
plt.ylabel('y')
```

Modeling I, the 1D case

Given a sample : (y_i, x_i) , for $i = 1, \dots, n$

Linear model or linear regression hypothesis assume :

$$y_i \approx \theta_0^* + \theta_1^* x_i$$

Model coefficients

- ▶ θ_0^* : intercept (unknown)
- ▶ θ_1^* : slope (unknown)

Rem: both parameters are unknown from the statistician

Data

- ▶ y is an **observation** or a variable to explain
- ▶ x is a **feature** or a covariate

Modeling II

Probabilistic model. Let us give a precise meaning to the sign \approx :

$$y_i = \theta_0^* + \theta_1^* x_i + \varepsilon_i,$$

$$\varepsilon_i \stackrel{i.i.d}{\sim} \varepsilon, \text{ for } i = 1, \dots, n$$

$$\mathbb{E}(\varepsilon) = 0$$

where i.i.d. means “independent and identically distributed”

Interpretation : $\varepsilon_i = y_i - \theta_0^* - \theta_1^* x_i$: represent the error between the theoretical model and the observations, represented by random variables ε_i centered (often referred to as **white noise**).

Rem: motivation for the random nature of the noise – measurement noise, transmission noise, in-population variability, etc.

Modeling III

$$y_i = \theta_0^* + \theta_1^* x_i + \varepsilon_i$$

We call


- **intercept** the scalar θ_0^* (■ ■ : *ordonnée à l'origine*)
- **slope** the scalar θ_1^* (■ ■ : *pente*)

Our **goal in the learning stage** is to estimate θ_0^* and θ_1^* (unknown) by $\hat{\theta}_0$ and $\hat{\theta}_1$ relying on observations (y_i, x_i) for $i = 1, \dots, n$


Rem: The “hat” notation is classical in statistics for referring to estimators

In **prediction time** $\hat{y}_i = \hat{\theta}_0 + \hat{\theta}_1 x_i$


Least squares : visualization




Least squares : visualization



(Total) Least squares : visualization



(Total) Least squares : visualization



Learning : mathematical formulation of Least squares

The **least squares** estimator is defined as :

$$(\hat{\theta}_0, \hat{\theta}_1) \in \operatorname{argmin}_{(\theta_0, \theta_1) \in \mathbb{R}^2} \sum_{i=1}^n (y_i - \theta_0 - \theta_1 x_i)^2$$

- ▶ Differentiate between θ^* , θ and $\hat{\theta}$!!!!
- ▶ it is also referred to as “ordinary least squares” (OLS)
- ▶ an original motivation for the squares is computational : first order conditions only require solving a linear system
- ▶ a solution always exists : minimizing a **coercive** continuous function
(coercive : $\lim_{\|x\| \rightarrow +\infty} f(x) = +\infty$)

Rem: write « $\in \operatorname{argmin}$ » as long as you do not know if the solution is unique

Least square authorship (controversial)



FIGURE – Adrien-Marie Legendre and Carl Friedrich Gauss

Historical / robust detour

The **least absolute deviation** (LAD) estimator reads :

$$(\hat{\theta}_0, \hat{\theta}_1) \in \operatorname{argmin}_{(\theta_0, \theta_1) \in \mathbb{R}^2} \sum_{i=1}^n |y_i - \theta_0 - \theta_1 x_i|$$

Rem: hard to compute without computer ; requires an optimization solver for non-smooth function (or a Linear Programming solver)

Rem: more robust to outliers (■ ■ : *données aberrantes*)

Least absolute deviation authorship





FIGURE – Ruđer Josip Bošković and Pierre-Simon de Laplace

Existence and uniqueness of the solution

Existence of a Local minimum : first order condition


Fermat's rule Theorem If f is differentiable, then at a local minimum x^* the gradient of f vanishes at x^* , *i.e.* $\nabla f(x^*) = 0$.



Existence and uniqueness of the solution

Existence of a Local minimum : first order condition

Fermat's rule Theorem If f is differentiable, then at a local minimum x^* the gradient of f vanishes at x^* , *i.e.* $\nabla f(x^*) = 0$.




Rem: sufficient condition when f is strongly convex !

Existence and uniqueness of the solution

Existence of a Local minimum : first order condition

Fermat's rule Theorem If f is differentiable, then at a local minimum x^* the gradient of f vanishes at x^* , *i.e.* $\nabla f(x^*) = 0$.



Rem: sufficient condition when f is strongly convex !

The Hessian Matrix and Gradients

The **gradient** ∇f is a vector of first-order partial derivatives :

$$\nabla f = \begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \frac{\partial f}{\partial x_2} \\ \vdots \\ \frac{\partial f}{\partial x_n} \end{bmatrix}$$

The **Hessian Matrix** \mathbf{H} of f is a square matrix of second-order partial derivatives :

$$\mathbf{H} = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \cdots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}$$

The minimizer is unique $\Leftrightarrow f$ is strictly convex

f is quadratic $\implies f$ is convex.

$f(\hat{\boldsymbol{\theta}})$ strictly convex $\Leftrightarrow \nabla^2 f(\hat{\boldsymbol{\theta}})$ positive definite $\Leftrightarrow \det(\nabla^2 f(\hat{\boldsymbol{\theta}})) > 0$

Back to least squares

$$\hat{\boldsymbol{\theta}} = (\hat{\theta}_0, \hat{\theta}_1) \in \operatorname{argmin}_{(\theta_0, \theta_1) \in \mathbb{R}^2} \frac{1}{2} \sum_{i=1}^n (y_i - \theta_0 - \theta_1 x_i)^2$$

For least squares, minimize the function of two variables :

$$f(\theta_0, \theta_1) = f(\boldsymbol{\theta}) = \frac{1}{2} \sum_{i=1}^n (y_i - \theta_0 - \theta_1 x_i)^2$$

First order condition / Fermat's rule :

$$\begin{cases} \frac{\partial f}{\partial \theta_0}(\hat{\boldsymbol{\theta}}) = \sum_{i=1}^n (y_i - \hat{\theta}_0 - \hat{\theta}_1 x_i) = 0 \\ \frac{\partial f}{\partial \theta_1}(\hat{\boldsymbol{\theta}}) = \sum_{i=1}^n (y_i - \hat{\theta}_0 - \hat{\theta}_1 x_i) x_i = 0 \end{cases}$$

Calculus continued

$$\text{Usual mean notation : } \bar{x}_n = \frac{1}{n} \sum_{i=1}^n x_i \text{ and } \bar{y}_n = \frac{1}{n} \sum_{i=1}^n y_i$$

With that, Fermat's rule states (dividing by n) :

$$\begin{cases} \frac{\partial f}{\partial \theta_0}(\hat{\boldsymbol{\theta}}) = \sum_{i=1}^n (y_i - \hat{\theta}_0 - \hat{\theta}_1 x_i) = 0 \\ \frac{\partial f}{\partial \theta_1}(\hat{\boldsymbol{\theta}}) = \sum_{i=1}^n (y_i - \hat{\theta}_0 - \hat{\theta}_1 x_i) x_i = 0 \end{cases}$$


\Leftrightarrow

$$\begin{cases} \hat{\theta}_0 = \bar{y}_n - \hat{\theta}_1 \bar{x}_n & \text{(CNO1)} \\ \hat{\theta}_1 = \frac{\sum_{i=1}^n (x_i - \bar{x}_n)(y_i - \bar{y}_n)}{\sum_{i=1}^n (x_i - \bar{x}_n)^2} & \text{(CNO2)} \end{cases}$$

Exo : Show that the solution to the OLS is unique iff $Var(x) \neq 0$

Center of gravity and interpretation

$$(\text{CNO1}) \Leftrightarrow (\bar{x}_n, \bar{y}_n) \in \{(x, y) \in \mathbb{R}^2 : y = \hat{\theta}_0 + \hat{\theta}_1 x\}$$



► $\overline{speed} = 15.4$

► $\overline{dist} = 42.98$

► $\hat{\theta}_0 = -17.579095$ intercept (negative!)

► $\hat{\theta}_1 = 3.932409$ slope

Physical interpretation : the cloud of points' center of gravity belongs to the (estimated) regression line

Vector formulation

Notation : $\mathbf{x} = (x_1, \dots, x_n)^\top$ and $\mathbf{y} = (y_1, \dots, y_n)^\top$

$$(\text{CNO2}) \Leftrightarrow \hat{\theta}_1 = \frac{\sum_{i=1}^n (x_i - \bar{x}_n)(y_i - \bar{y}_n)}{\sum_{i=1}^n (x_i - \bar{x}_n)^2}$$

$$(\text{CNO2}) \Leftrightarrow \hat{\theta}_1 = \text{corr}_n(\mathbf{x}, \mathbf{y}) \cdot \frac{\sqrt{\text{var}_n(\mathbf{y})}}{\sqrt{\text{var}_n(\mathbf{x})}}$$

where $\text{corr}_n(\mathbf{x}, \mathbf{y}) = \frac{\frac{1}{n} \sum_{i=1}^n (x_i - \bar{x}_n)(y_i - \bar{y}_n)}{\sqrt{\text{var}_n(\mathbf{x})} \sqrt{\text{var}_n(\mathbf{y})}}$


and $\text{var}_n(\mathbf{z}) = \frac{1}{n} \sum_{i=1}^n (z_i - \bar{z}_n)^2$ (for any $\mathbf{z} = (z_1, \dots, z_n)^\top$)

respectively **empirical correlation**, **empirical variances**

cars example

Braking distance for cars as a function of the speed

Line slope : $\text{corr}_n(\mathbf{x}, \mathbf{y}) \cdot \frac{\sqrt{\text{var}_n(\mathbf{y})}}{\sqrt{\text{var}_n(\mathbf{x})}} = 3.932409$.



Centering

Centered model :


$$\text{Write for any } i = 1, \dots, n : \begin{cases} x'_i = x_i - \bar{x}_n \\ y'_i = y_i - \bar{y}_n \end{cases} \Leftrightarrow \begin{cases} \mathbf{x}' = \mathbf{x} - \bar{x}_n \mathbf{1}_n \\ \mathbf{y}' = \mathbf{y} - \bar{y}_n \mathbf{1}_n \end{cases}$$

and $\mathbf{1}_n = (1, \dots, 1)^\top \in \mathbb{R}^n$, then solving the OLS with $(\mathbf{x}', \mathbf{y}')$ leads to

$$\begin{cases} \hat{\theta}'_0 = 0 \\ \hat{\theta}'_1 = \frac{\frac{1}{n} \sum_{i=1}^n x'_i y'_i}{\frac{1}{n} \sum_{i=1}^n x'^2_i} \end{cases}$$

Rem: equivalent to choosing the cloud of points' center of mass as origin, *i.e.*
 $(\bar{x}'_n, \bar{y}'_n) = (0, 0)$

Centering (II)




Centering and interpretation

Consider the coefficient $\hat{\theta}'_1$ ($\hat{\theta}'_0 = 0$) for centered points \mathbf{y}', \mathbf{x}' , then :


$$\hat{\theta}'_1 \in \operatorname{argmin}_{\theta_1} \sum_{i=1}^n (y'_i - \theta_1 x'_i)^2 = \operatorname{argmin}_{\theta_1} \sum_{i=1}^n x_i'^2 \left(\frac{y'_i}{x'_i} - \theta_1 \right)^2$$

Interpretation : $\hat{\theta}'_1$ is a weighted average of the slopes $\frac{y'_i}{x'_i}$


$$\hat{\theta}'_1 = \frac{\sum_{i=1}^n x_i'^2 \frac{y'_i}{x'_i}}{\sum_{j=1}^n x_j'^2}$$

Influence of extreme points : weights proportional to $x_i'^2$; connected to the **leverage** ( : *levier*) effect


Extreme points – leverage effect



Extreme points – leverage effect



Extreme points – leverage effect (II)



Centering + scaling (standardization)

Centered-scaled model :

$$\forall i = 1, \dots, n : \begin{cases} x_i'' = (x_i - \bar{x}_n) / \sqrt{\text{var}_n(\mathbf{x})} \\ y_i'' = (y_i - \bar{y}_n) / \sqrt{\text{var}_n(\mathbf{y})} \end{cases} \Leftrightarrow \begin{cases} \mathbf{x}'' = \frac{\mathbf{x} - \bar{x}_n \mathbf{1}_n}{\sqrt{\text{var}_n(\mathbf{x})}} \\ \mathbf{y}'' = \frac{\mathbf{y} - \bar{y}_n \mathbf{1}_n}{\sqrt{\text{var}_n(\mathbf{y})}} \end{cases}$$

Solving OLS with $(\mathbf{x}'', \mathbf{y}'')$ then

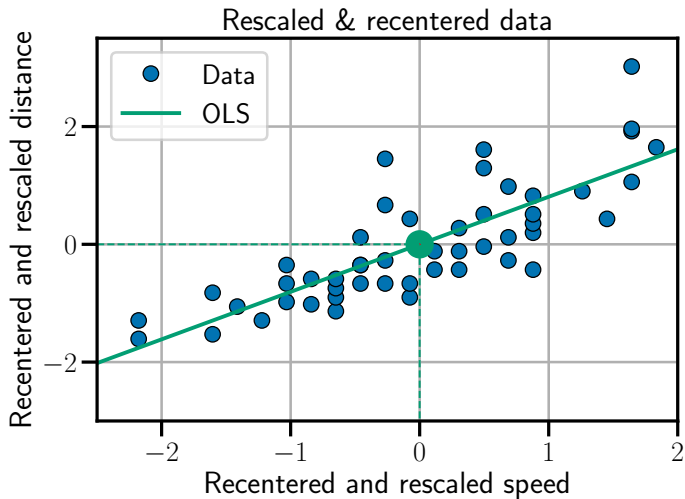
$$\begin{cases} \hat{\theta}_0'' = 0 \\ \hat{\theta}_1'' = \frac{1}{n} \sum_{i=1}^n x_i'' y_i'' \end{cases}$$

Rem: equivalent to choosing the points cloud center of mass as origin and normalize \mathbf{x} and \mathbf{y} to have unit **empirical norm** $\|\cdot\|_n$:

$$\|\mathbf{x}''\|_n^2 = \frac{1}{n} \sum_{i=1}^n (x_i'')^2 = 1$$

$$\|\mathbf{y}''\|_n^2 = \frac{1}{n} \sum_{i=1}^n (y_i'')^2 = 1$$

Centering + scaling



When/why preprocessing?

Centering \mathbf{y} or using an intercept (or adding a constant feature) is equivalent

Rem: for sparse (■ ■ : *creux*) cases centering \mathbf{y} adding a constant feature could be preferred

Scaling features is important :

- ▶ if you want to interpret the coefficients' amplitude in regression (better solution : t-tests)
- ▶ if you want to penalize or regularize coefficients (*c.f.* Lasso, Ridge, etc.) a single scale is needed
- ▶ for computing reasons (*e.g.* store scaling to improve efficiency, etc.)

Rem: in practice centering/scaling is useful for **estimation** not so much for **prediction** (see next courses)

What happens with the logarithm scaling?

Centering with Python

Use centering classes from `sklearn`, see preprocessing :

<http://scikit-learn.org/stable/modules/preprocessing.html>

```
from sklearn import preprocessing

scaler = preprocessing.StandardScaler().fit(X)

print(np.isclose(scaler.mean_, np.mean(X)))

print(np.array_equal(scaler.std_, np.std(X)))

print(np.array_equal(scaler.transform(X),
                    (X - np.mean(X)) / np.std(X)))

print(np.array_equal(scaler.transform([26]),
                    (26 - np.mean(X)) / np.std(X)))
```

Rem: most valuable with pipeline

<http://scikit-learn.org/stable/modules/pipeline.html>

Prediction

We call **prediction** function the function that associates an estimation of the variable of interest to a new sample. For least squares the prediction is given by :

$$\text{pred}(x_{n+1}) = \hat{\theta}_0 + \hat{\theta}_1 x_{n+1}$$

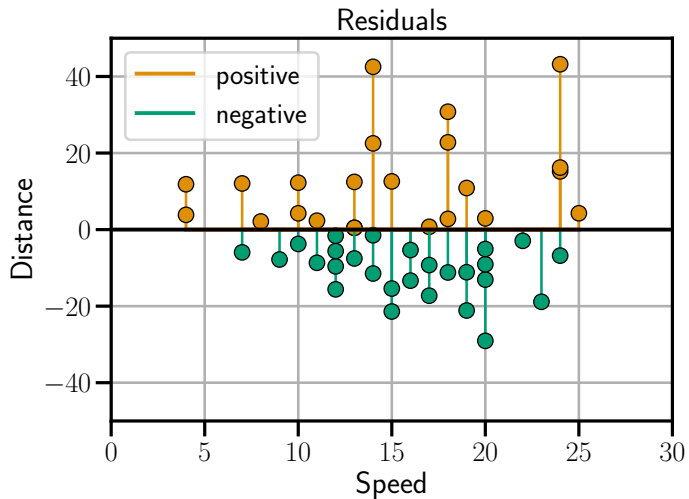
Rem: often written \hat{y}_{n+1} (implicit dependence on x_{n+1})

The **residual** : difference between observations and predicted values


$$\epsilon_i = y_i - \text{pred}(x_i) = y_i - \hat{y}_i = y_i - (\hat{\theta}_0 + \hat{\theta}_1 x_i)$$

Rem: observable estimate of the unobservable statistical error

Residuals (on cars)



Residual histograms



Least squares motivation

- ▶ Computing advantage : computationally heavy methods avoided before computers (*e.g.* iterative methods)
- ▶ Theoretical advantage : least square analysis easy under simple hypothesis
- ▶ Interpretability : how much does the regressor increase with the features


Example : under additive white Gaussian noise assumption *i.e.*, $\varepsilon \sim \mathcal{N}(0, \sigma^2)$ the maximum likelihood is equivalent to solving least squares to estimate (θ_0^*, θ_1^*)

Rem: for another noise model and/or to limit outliers influence one can solve (see *e.g.* `QuantReg` in `statsmodels`)

$$\hat{\theta} = (\hat{\theta}_0, \hat{\theta}_1) \in \operatorname{argmin}_{(\theta_0, \theta_1) \in \mathbb{R}^2} \sum_{i=1}^n |y_i - \theta_0 - \theta_1 x_i|$$

Discussion : toward multivariate cases

Physical laws (or your driving school memories) would lead to rather pick a **quadratic** model instead of a **linear** one : the OLS can be applied by choosing x_i^2 as features instead of x_i :




Web sites and books to go further

- ▶ Datascience in general : Blog + videos by Jake Vanderplas
<http://jakevdp.github.io/>
Homework for next lesson : watch the following videos <http://jakevdp.github.io/blog/2017/03/03/reproducible-data-analysis-in-jupyter/>
- ▶ A few [notebooks](#) of OLS with statsmodels
- ▶ [McKinney \(2012\)](#) about Python for statistics
- ▶ [Lejeune \(2010\)](#) about linear models (in French)
- ▶ Regression course by [B. Delyon](#) (in French, more technical)


Toward multivariate models

Tree volume as a function of height / girth (■ : *circonférence*)



Toward multivariate models

Tree volume as a function of height / girth (■ ■ : *circonférence*)



Python commands

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression

# Generate example data
...

# Fit linear regression model
model = LinearRegression()
model.fit(X, y)
```

Model

One observes p features $(\mathbf{x}_1, \dots, \mathbf{x}_p)$. Model in dimension p

$$y_i = \theta_0^* + \sum_{j=1}^p \theta_j^* x_{i,j} + \varepsilon_i$$

$$\varepsilon_i \stackrel{i.i.d}{\sim} \varepsilon, \text{ pour } i = 1, \dots, n$$

$$\mathbb{E}[\varepsilon] = 0$$

Rem: we assume (frequentist point of view) there exists a “true” parameter

$$\boldsymbol{\theta}^* = (\theta_0^*, \dots, \theta_p^*)^\top \in \mathbb{R}^{p+1}$$

Dimension p

Matrix model

$$\underbrace{\begin{pmatrix} y_1 \\ \vdots \\ y_n \end{pmatrix}}_{\mathbf{y}} = \underbrace{\begin{pmatrix} 1 & x_{1,1} & \dots & x_{1,p} \\ \vdots & \vdots & \ddots & \vdots \\ 1 & x_{n,1} & \dots & x_{n,p} \end{pmatrix}}_X \underbrace{\begin{pmatrix} \theta_0^* \\ \vdots \\ \theta_p^* \end{pmatrix}}_{\boldsymbol{\theta}^*} + \underbrace{\begin{pmatrix} \varepsilon_1 \\ \vdots \\ \varepsilon_n \end{pmatrix}}_{\boldsymbol{\epsilon}}$$

Equivalently : $\boxed{\mathbf{y} = X\boldsymbol{\theta}^* + \boldsymbol{\epsilon}}$ (1)

Column notation : $X = (\mathbf{x}_0, \mathbf{x}_1, \dots, \mathbf{x}_p)$ with $\mathbf{x}_0 = \mathbf{1}_n = \begin{pmatrix} 1 \\ \vdots \\ 1 \end{pmatrix}$.

Line notation : $X = \begin{pmatrix} x_1^\top \\ \vdots \\ x_n^\top \end{pmatrix} = (x_1, \dots, x_n)^\top$

Matrix Notation and L_2 Norm

Matrix notation is a powerful way to represent mathematical operations involving vectors and matrices.

The **Inner Product** (dot product) of two vectors \mathbf{u} and \mathbf{v} is defined as :

$$\langle \mathbf{u}, \mathbf{v} \rangle = \sum_{i=1}^n u_i v_i = \mathbf{u} \cdot \mathbf{v}^T$$

Let \mathbf{A} be an $m \times n$ matrix and \mathbf{B} be an $n \times p$ matrix. The **matrix product** $\mathbf{C} = \mathbf{AB}$ is an $m \times p$ matrix with elements :

$$c_{ij} = \sum_{k=1}^n a_{ik} b_{kj}$$

The L_2 **Norm** (Euclidean norm) of a vector \mathbf{v} is defined as :

$$\|\mathbf{v}\|_2 = \sqrt{\sum_{i=1}^n v_i^2}$$

Matrix notation simplifies operations and equations involving vectors and matrices.

Vocabulary

$$\mathbf{y} = X\boldsymbol{\theta}^* + \boldsymbol{\epsilon}$$

- ▶ $\mathbf{y} \in \mathbb{R}^n$: observations vector
- ▶ $X \in \mathbb{R}^{n \times (p+1)}$: **design** matrix (with features as columns and a first column of 1s)
- ▶ $\tilde{X} \in \mathbb{R}^{n \times (p)}$: **reduced design** matrix (with features as columns and NO column of ones)
- ▶ $\boldsymbol{\theta}^* \in \mathbb{R}^{p+1}$: (unknown) **true** parameter to be estimated
- ▶ $\boldsymbol{\epsilon} \in \mathbb{R}^n$: noise vector

(Ordinary) Least squares

A least square estimator is any solution of the following problem :

$$\hat{\boldsymbol{\theta}} \in \operatorname{argmin}_{\boldsymbol{\theta} \in \mathbb{R}^{p+1}} \left(\|\mathbf{y} - X\boldsymbol{\theta}\|_2^2 \right)$$

$$\hat{\boldsymbol{\theta}} \in \operatorname{argmin}_{\boldsymbol{\theta} \in \mathbb{R}^{p+1}} \sum_{i=1}^n \left[y_i - \left(\theta_0 + \sum_{j=1}^p \theta_j x_{i,j} \right) \right]^2$$

$$\hat{\boldsymbol{\theta}} \in \operatorname{argmin}_{\boldsymbol{\theta} \in \mathbb{R}^{p+1}} \sum_{i=1}^n [y_i - \langle x_i, \boldsymbol{\theta} \rangle]^2$$

- Does the solution exist ? A solution always exists, as we are minimizing a coercive continuous function (**coercive** : $\lim_{\|x\| \rightarrow +\infty} f(x) = +\infty$)
- Is the solution unique ? not guaranteed

Exo how do we make the prediction ?

Row / column interpretation

Row interpretation

Let $\tilde{x}_1^\top, \dots, \tilde{x}_{p+1}^\top$ be the rows of X . The residuals are $r_i = \tilde{x}_i \boldsymbol{\theta} - y_i$ and the OLS is equivalent to minimizing the sum of squares residuals

Column interpretation

Let x_1, \dots, x_{p+1} be the columns of X . Then $\|\mathbf{y} - X\boldsymbol{\theta}\|_2^2 = \|(\theta_0 x_0, \dots, \theta_p x_p) - \mathbf{y}\|_2^2$, so OLS is to find a linear combination of columns of X that is closest to \mathbf{y} .

Vocabulary (and abuse of terms)

We call **Gram matrix** the matrix

$$X^\top X$$

whose general term is $[X^\top X]_{i,j} = \langle \mathbf{x}_i, \mathbf{x}_j \rangle$

If the design matrix X is centered and scaled, the Gram matrix is proportional to the correlation between columns. $X^\top X$ is often referred to as the feature correlation matrix

Rem: when columns are scaled such that $\forall j \in \llbracket 0, p \rrbracket, \|\mathbf{x}_j\|^2 = n$, the Gramian diagonal is (n, \dots, n)

The vector $X^\top \mathbf{y} = \begin{pmatrix} \langle \mathbf{x}_0, \mathbf{y} \rangle \\ \vdots \\ \langle \mathbf{x}_p, \mathbf{y} \rangle \end{pmatrix}$ represents the correlation between the observations and the features

Hilbert projection theorem (HPT)

Let $C \subset \mathbb{R}^d, Y \in \mathbb{R}^d$. Let $\hat{z} = \arg \min_{z \in C} \|Y - z\|_2^2$. Then \hat{z} always exists and is given by

$$\langle Y - \hat{z}, z \rangle = 0 \quad \forall z \in C$$

Hilbert projection theorem (HPT) and application to OLS

$$\hat{\boldsymbol{\theta}} \in \operatorname{argmin}_{\boldsymbol{\theta} \in \mathbb{R}^{p+1}} \|\mathbf{y} - X\boldsymbol{\theta}\|_2^2$$

Note $\operatorname{col}(X) = \operatorname{span}([x_0, \dots, x_p]) = \sum_{j=0}^p x_j \boldsymbol{\theta}_j = X\boldsymbol{\theta}$ OLS :

$$\widehat{W} \in \operatorname{argmin}_{W \in \operatorname{col}(X)} (\|\mathbf{y} - W\|_2^2)$$

$$\langle \mathbf{y} - \widehat{W}, W \rangle = 0$$

$$(\mathbf{y} - \widehat{W})^\top W = 0$$

$$(\mathbf{y} - \widehat{W})^\top X\boldsymbol{\theta} = 0$$

$$(\mathbf{y} - \widehat{W})^\top X = 0 \tag{2}$$

$$(\mathbf{y} - X\hat{\boldsymbol{\theta}})^\top X = 0$$

$$X^\top (\mathbf{y} - X\hat{\boldsymbol{\theta}}) = 0$$

$$X^\top X\hat{\boldsymbol{\theta}} = X^\top \mathbf{y}$$

OLS normal equations

The solution to the OLS problem is given by the solution to the normal equation

$$\text{Normal equation : } \boxed{X^{\top} X \hat{\boldsymbol{\theta}} = X^{\top} \mathbf{y}}$$

As a consequence,

- ▶ a solution always exists.
- ▶ its unique if the solution to the normal equations is unique

Hilbert projection theorem




FIGURE — Source : Wikipedia

Least squares and uniqueness

Let $\hat{\boldsymbol{\theta}}$ be a solution of $\boxed{X^\top X \hat{\boldsymbol{\theta}} = X^\top \mathbf{y}}$

Non uniqueness : happens for non trivial kernel, *i.e.* when $\ker(X) = \{\boldsymbol{\theta} \in \mathbb{R}^{p+1} : X\boldsymbol{\theta} = 0\} \neq \{0\}$

Assume $\boldsymbol{\theta}_K \in \ker(X)$ with $\boldsymbol{\theta}_K \neq 0$, then

$$X(\hat{\boldsymbol{\theta}} + \boldsymbol{\theta}_K) = X\hat{\boldsymbol{\theta}}$$


$$\text{and then } (X^\top X)(\hat{\boldsymbol{\theta}} + \boldsymbol{\theta}_K) = X^\top \mathbf{y}$$

Conclusion : the set of least squares solutions is an affine sub-space

$$\boxed{\hat{\boldsymbol{\theta}} + \ker(X)}$$


Optimization in \mathbb{R}^d

Convex case, $f(\boldsymbol{\theta}) = \|\mathbf{y} - X\boldsymbol{\theta}\|_2^2$, where the set of minimizers is non-unique :




Optimization in \mathbb{R}^d

Convex case, $f(\boldsymbol{\theta}) = \|\mathbf{y} - X\boldsymbol{\theta}\|_2^2$, where the set of minimizers is non-unique :




Optimization in \mathbb{R}^d

Convex case, $f(\boldsymbol{\theta}) = \|\mathbf{y} - X\boldsymbol{\theta}\|_2^2$, where the set of minimizers is non-unique :




Optimization in \mathbb{R}^d

Convex case, $f(\boldsymbol{\theta}) = \|\mathbf{y} - X\boldsymbol{\theta}\|_2^2$, where the set of minimizers is non-unique :



Optimization in \mathbb{R}^d

Convex case, $f(\boldsymbol{\theta}) = \|\mathbf{y} - X\boldsymbol{\theta}\|_2^2$, where the set of minimizers is non-unique :



Non uniqueness : single feature case

Reminder :

$$X = \begin{pmatrix} 1 & x_1 \\ \vdots & \vdots \\ 1 & x_n \end{pmatrix}$$

If $\ker(X) = \{\boldsymbol{\theta} \in \mathbb{R}^2 : X\boldsymbol{\theta} = 0\} \neq \{0\}$ there exists $(\theta_0, \theta_1) \neq (0, 0)$:

$$\begin{cases} \theta_0 + \theta_1 x_1 & = 0 \\ \vdots & \vdots & = \vdots \\ \theta_0 + \theta_1 x_n & = 0 \end{cases} \quad (\star)$$

1. If $\theta_1 = 0$: $(\star) \Rightarrow \theta_0 = 0$, so $(\theta_0, \theta_1) = (0, 0)$, **contradiction**

2. If $\theta_1 \neq 0$:

2.1 If $\forall i, x_i = 0$ then $X = (\mathbf{1}_n, 0)$ and $\theta_0 = 0$

2.2 Otherwise there exists $x_{i_0} \neq 0$ and $\forall i, x_i = -\theta_0/\theta_1 = x_{i_0}$, *i.e.* $X = [\mathbf{1}_n \quad x_{i_0} \cdot \mathbf{1}_n]$

Interpretation : $\mathbf{x}_1 \propto \mathbf{1}_n$, *i.e.* \mathbf{x}_1 is constant

Interpretation for multivariate cases

Reminder : we write $X = (\mathbf{1}_n, \mathbf{x}_1, \dots, \mathbf{x}_p)$, the features being column-wise (each are of length n)

The property $\ker(X) = \{\boldsymbol{\theta} \in \mathbb{R}^{p+1} : X\boldsymbol{\theta} = 0\} \neq \{0\}$ means that there exists a linear dependence between the features $\mathbf{1}_n, \mathbf{x}_1, \dots, \mathbf{x}_p$,

Reformulation : $\exists \boldsymbol{\theta} = (\theta_0, \dots, \theta_p)^\top \in \mathbb{R}^{p+1} \setminus \{0\}$ s.t.

$$\theta_0 \mathbf{1}_n + \sum_{j=1}^p \theta_j \mathbf{x}_j = 0$$

Algebra reminder

Rank of a matrix : $\text{rank}(X) = \dim(\text{span}(\mathbf{1}_n, \mathbf{x}_1, \dots, \mathbf{x}_p))$; $\text{span}(\cdot)$: the space generated by \cdot .

Property : $\text{rank}(X) = \text{rank}(X^\top)$

Rank–nullity theorem :

- ▶ $\text{rank}(X) + \dim(\ker(X)) = p + 1$
- ▶ $\text{rank}(X^\top) + \dim(\ker(X^\top)) = n$

Property : $\boxed{\text{rank}(X) \leq \min(n, p + 1)}$

See [Golub and Van Loan \(1996\)](#) for details

Algebra reminder (continued)

Matrix inversion : A square matrix $A \in \mathbb{R}^{m \times m}$ is invertible

- ▶ if and only if its kernel is trivial : $\ker(A) = \{0\}$
- ▶ if and only if it is full rank $\text{rank}(A) = m$

OLS is unique iff $X^\top X$ is invertible

$$\Leftrightarrow \ker(X^\top X) = \{0\}$$

$$\Leftrightarrow \ker(X) = \{0\}$$

$$\Leftrightarrow X \text{ has full rank}$$

Exo: $\ker(X) = \ker(X^\top X)$

Prediction

$$\text{Prediction vector : } \hat{\mathbf{y}} = X\hat{\boldsymbol{\theta}}$$

Rem: $\hat{\mathbf{y}}$ depends linearly on the observation vector \mathbf{y}

Rem: an **orthogonal projector** is a matrix H such that

1. H is symmetric : $H^\top = H$
2. H is idempotent : $H^2 = H$

Proposition Writing H_X the orthogonal projector onto the space span by the columns of X , one gets $\hat{\mathbf{y}} = H_X \mathbf{y}$

If X is full (column) rank, then $H_X = X(X^\top X)^{-1}X^\top$ is called the **hat matrix**

Exo: Show that H_X is an orthogonal projector

Prediction (continued)

If a new observation $x_{n+1} = (x_{n+1,1}, \dots, x_{n+1,p})$ is provided, the associated prediction is :

$$\hat{y}_{n+1} = \langle \hat{\boldsymbol{\theta}}, (1, x_{n+1,1}, \dots, x_{n+1,p})^\top \rangle$$

$$\hat{y}_{n+1} = \hat{\theta}_0 + \sum_{j=1}^p \hat{\theta}_j x_{n+1,j}$$

Rem: the normal equation ensures **equi-correlation** between observations and features :

$$\begin{aligned} (X^\top X) \hat{\boldsymbol{\theta}} &= X^\top \mathbf{y} \Leftrightarrow X^\top \hat{\mathbf{y}} = X^\top \mathbf{y} \\ &\Leftrightarrow \begin{pmatrix} \langle \mathbf{x}_0, \hat{\mathbf{y}} \rangle \\ \vdots \\ \langle \mathbf{x}_p, \hat{\mathbf{y}} \rangle \end{pmatrix} = \begin{pmatrix} \langle \mathbf{x}_0, \mathbf{y} \rangle \\ \vdots \\ \langle \mathbf{x}_p, \mathbf{y} \rangle \end{pmatrix} \end{aligned}$$

Residuals and normal equation

$$\text{Residual(s) : } \hat{\boldsymbol{\varepsilon}} = \mathbf{y} - \hat{\mathbf{y}} = \mathbf{y} - X\hat{\boldsymbol{\theta}} = (\text{Id}_n - H_X)\mathbf{y}$$

Proposition

$$\begin{aligned} \langle \hat{\boldsymbol{\varepsilon}}, X \rangle &= 0_n \\ \langle \hat{\boldsymbol{\varepsilon}}, \hat{\mathbf{y}} \rangle &= 0 \\ \langle \hat{\boldsymbol{\varepsilon}}, \bar{\mathbf{y}}\mathbf{1}_n \rangle &= 0 \end{aligned} \tag{3}$$

Rem: The Normal equation is $(X^\top X)\hat{\boldsymbol{\theta}} = X^\top \mathbf{y}$. It follows that $X^\top(X\hat{\boldsymbol{\theta}} - \mathbf{y}) = 0 \Leftrightarrow X^\top \hat{\boldsymbol{\varepsilon}} = 0 \Leftrightarrow \hat{\boldsymbol{\varepsilon}}^\top X = 0$

With $X = (\mathbf{1}_n, \mathbf{x}_1, \dots, \mathbf{x}_p)$, this can be rewritten

$$\forall j = 1, \dots, p : \langle \hat{\boldsymbol{\varepsilon}}, \mathbf{x}_j \rangle = 0 \text{ and } \bar{r}_n = 0$$

Interpretation : (1,2) residuals are \perp to features and (3) $\hat{\boldsymbol{\varepsilon}}$ is centered ($\sum \hat{\varepsilon}_i = 0$)

How good is our model ? RSS and the determination coefficient R^2

The ratio of the variation explained by the model and the total variation of the data $R^2 = \frac{\|\hat{\mathbf{y}} - \bar{\mathbf{y}}\mathbf{1}_n\|^2}{\|\mathbf{y} - \bar{\mathbf{y}}\mathbf{1}_n\|^2}$ We can write also, by orthogonality :

$$\|\mathbf{y} - \bar{\mathbf{y}}\mathbf{1}_n\|^2 = \|\mathbf{y} - \hat{\mathbf{y}}\|^2 + \|\hat{\mathbf{y}} - \bar{\mathbf{y}}\mathbf{1}_n\|^2 \quad (4)$$

Reordering


$$\|\hat{\mathbf{y}} - \bar{\mathbf{y}}\mathbf{1}_n\|^2 = \|\mathbf{y} - \bar{\mathbf{y}}\mathbf{1}_n\|^2 - \|\mathbf{y} - \hat{\mathbf{y}}\|^2 \quad (5)$$

So

$$R^2 = 1 - \frac{\|\mathbf{y} - \hat{\mathbf{y}}\|^2}{\|\mathbf{y} - \bar{\mathbf{y}}\mathbf{1}_n\|^2} \quad (6)$$





Exo: Show that $0 \leq R^2 \leq 1$

Polynomial regression and overfitting



Source : sklearn

References I

-  B. Delyon.
Régression, 2015.
<https://perso.univ-rennes1.fr/bernard.delyon/regression.pdf>.
-  G. H. Golub and C. F. van Loan.
Matrix computations.
Johns Hopkins University Press, Baltimore, MD, third edition, 1996.
-  M. Lejeune.
Statistiques, la théorie et ses applications.
Springer, 2010.
-  W. McKinney.
Python for Data Analysis : Data Wrangling with Pandas, NumPy, and IPython.
O'Reilly Media, 2012.